

Benchmarking P2C for HPC using NPB

Joseph Anthony C. Hermocilla

December 23, 2016

Abstract

We report some benchmarking results of the Peak-Two Cloud(P2C) for High-Performance Computing(HPC) using the NAS Parallel Benchmarks(NPB).

1 Introduction

In order to evaluate the performance of the P2C[1] for HPC, NPB¹ version 3.3.1 was run on a 16-node MPI cluster provisioned using vcluster².

NPB consists of a set of programs that implements different computational approaches associated with Computational Fluid Dynamics(CFD). These programs represent the types of applications that are run in supercomputers and HPC clusters. In running the benchmark, classes A and B were used for each program. Different classes have different problem sizes and parameters which result to different measurements.

The programs were run using 1, 2, 4, 8, and 16 nodes. The completion time were recorded and plotted. The figure below shows the hosts where the VMs used in the cluster were instantiated.

2 Results

The following subsections show the completion time for each program in NPB tested using different number of nodes. Ideally, the completion time of the programs should decrease as the number of nodes is increased. However, the results show some fluctuations for some programs.

2.1 Conjugate Gradient

Performs irregular memory access and communication.




Figure 1:

¹<https://www.nas.nasa.gov/publications/npb.html>

²<http://srg.ics.uplb.edu.ph/projects/peak-two-cloud/peak-two-cloud-resources/deployinganmpiclustertestingvcluster>




Figure 2:

2.2 Embarassingly Parallel

An application with very minimal communication and synchronization to complete a task.




Figure 3:




Figure 4:

2.3 Fourier Transform

Performs all-to-all communication.




Figure 5:




Figure 6:

2.4 Integer Sort

Performs random memory access.




Figure 7:




Figure 8:

2.5 Lower-Upper Gauss-Seidel Solver

A pseudo application that solves a system of linear equations.




Figure 9:




Figure 10:

2.6 Multi-Grid

Solves differential equations using long- and short-distance communication and is memory intensive.




Figure 11:




Figure 12:

References

- [1] J. A. C. Hermocilla. P2c: Towards scientific computing on private clouds. In *Proceedings of the 12th National Conference on IT Education (NCITE 2014)*, pages 163–168. Philippine Society of Information Technology Educators, Oct. 2014.